

CS 325 Week 4-1

Intro to the relational model (DB Reading Packet 3)

Intro to the RELATIONAL model

- Some KEY definitions

Intro to the RELATIONAL model

- Some KEY definitions
- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation

Intro to the RELATIONAL model

- Some KEY definitions
- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - E.g., advisor = (advisor_id, advisor_lname, advisor_fname, advisor_phone)

Intro to the RELATIONAL model

- Some KEY definitions
- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - E.g., advisor = (advisor_id, advisor_lname, advisor_fname, advisor_phone)
 - (it MUST have at least ONE such SET of attributes!)

Intro to the RELATIONAL model

- Some KEY definitions
- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - E.g., advisor = (advisor_id, advisor_lname, advisor_fname, advisor_phone)
 - (it MUST have at least ONE such SET of attributes!)
 - consider this relation:
 - course_stats = (stu_id, class_id, class_grade, completion_date)
 - So, what may be our superkey in this relation?

Intro to the RELATIONAL model

- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - (it MUST have at least ONE such SET of attributes!)
 - consider this relation:
 - `course_stats = (stu_id, class_id, class_grade, completion_date)`
 - it is the case that `(stu_id, class_id)` is a superkey of `course_stats`
 - it is also the case the `(stu_id, class_id, class_grade)` is a superkey of `course_stats`

Intro to the RELATIONAL model

- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - (it MUST have at least ONE such SET of attributes!)
 - consider this relation:
 - `course_stats = (stu_id, class_id, class_grade, completion_date)`
 - it is the case that `(stu_id, class_id)` is a superkey of `course_stats`
 - it is also the case the `(stu_id, class_id, class_grade)` is a superkey of `course_stats`
 - Is `stu_id` alone a superkey? Is `Class_id` alone a superkey?

Intro to the RELATIONAL model

- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - (it MUST have at least ONE such SET of attributes!)
 - consider this relation:
 - `course_stats = (stu_id, class_id, class_grade, completion_date)`
 - it is the case that `(stu_id, class_id)` is a superkey of `course_stats`, it is also the case the `(stu_id, class_id, class_grade)` is a superkey of `course_stats`
 - Is the combination of `class_id` and `class_grade` a superkey? Is the combination of `class_id` and `completion_date` a superkey?

Intro to the RELATIONAL model

- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - (it MUST have at least ONE such SET of attributes!)
- PLEASE NOTE: the functional dependencies and superkeys of relations representing some setting or scenario are determined by that scenario, and the users' model of that scenario;

Intro to the RELATIONAL model

- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
 - a relation can have multiple superkeys!
 - (it MUST have at least ONE such SET of attributes!)
- PLEASE NOTE: the functional dependencies and superkeys of relations representing some setting or scenario are determined by that scenario, and the users' model of that scenario;
- And business rules - day-to-day rules in some setting or scenario, and such business rules can affect functional dependencies and superkeys;

Intro to the RELATIONAL model

- PLEASE NOTE: the functional dependencies and superkeys of relations representing some setting or scenario are determined by that scenario, and the users' model of that scenario;
- And business rules - day-to-day rules in some setting or scenario, and such business rules can affect functional dependencies and superkeys;
- So if you have a question about whether some relationship between attributes IS a functional dependency -- you should ask the users!!!

Question

Consider the relation:

$\text{dept}(\text{dept_name}, \text{DEPT_NUM}, \text{dept_loc})$.

Note: in this scenario, some departments share an office.

Is the attribute-set (dept_loc) also a **superkey** for this relation?

- 1) Yes
- 2) No

Intro to the RELATIONAL model

- Some KEY definitions
- A SUPERKEY or LOGICAL KEY is any set of one or more attributes that uniquely determine a tuple in a relation
- e.g., course_stats = (stu_id, class_id, class_grade, completion_date)
- SO -- a superkey may be a bit "more" than is needed to determine a tuple in a relation;

Intro to the RELATIONAL model

- Some KEY definitions
- e.g., course_stats = (stu_id, class_id, class_grade, completion_date)
- SO -- a superkey may be a bit "more" than is needed to determine a tuple in a relation;
- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey

Intro to the RELATIONAL model

- SO -- a superkey may be a bit "more" than is needed to determine a tuple in a relation;
- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - Consider relation advisor = (advisor_id, advisor_lname, advisor_fname)

Intro to the RELATIONAL model

- Another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - Consider relation advisor = (advisor_id, advisor_lname, advisor_fname)
 - for advisor
 - (advisor_id, advisor_lname, advisor_fname) is a superkey but not a minimal key, because subset (advisor_id) is also a superkey
 - BUT (advisor_id) IS a minimal key, because no subset of it is ALSO a superkey.

Intro to the RELATIONAL model

- Another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - Consider relation advisor = (advisor_id, advisor_lname, advisor_fname)
 - for course_stats,
 - (stu_id, class_id, class_grade) is a superkey but not a minimal key, because subset (stu_id, class_id) is also a superkey BUT (stu_id, class_id) IS a minimal key, because no subset of it is ALSO a superkey.

Intro to the RELATIONAL model

- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - a minimal key is also called a CANDIDATE key

Intro to the RELATIONAL model

- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - a minimal key is also called a CANDIDATE key
 - well...
 - a relation CAN have multiple minimal/candidate keys...!
 - student = (stu_id, stu_ssn, stu_lname, stu_fname, stu_email)
 - ...then (stu_id) is a minimal key
 - but (stu_ssn) is also a minimal key

Intro to the RELATIONAL model

- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - a minimal key is also called a CANDIDATE key
 - the primary key for a relation is the minimal/candidate key that we CHOOSE (from amongst its possible minimal/candidate keys) to serve as its primary key;

Intro to the RELATIONAL model

- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - a minimal key is also called a CANDIDATE key
 - the primary key for a relation is the minimal/candidate key that we CHOOSE (from amongst its possible minimal/candidate keys) to serve as its primary key;
 - student = (stu_id, stu_ssn, stu_lname, stu_fname, stu_email)

Intro to the RELATIONAL model

- another definition is then:
 - a MINIMAL key is a superkey in which NO (proper) SUBSET of the attributes is ALSO a superkey
 - a minimal key is also called a CANDIDATE key
 - the primary key for a relation is the minimal/candidate key that we CHOOSE (from amongst its possible minimal/candidate keys) to serve as its primary key;
 - artworks = (art_id, art_title, artist_name, year_created, price)

Question

Again consider the relation:

$\text{dept}(\text{dept_name}, \text{DEPT_NUM}, \text{dept_loc})$.

And, again, remember that, in this scenario, some departments share an office.

Is the attribute-set $(\text{dept_num}, \text{dept_loc})$ a **candidate** key for this relation?

- a) Yes
- b) No

Intro to the RELATIONAL model

Another KEY point

- do you see that, given these definitions, it MUST be the case that, for a relation scheme (e.g., $R(A_1, A_2, a_3, a_4, a_5)$), the primary key attributes MUST functionally determine the remaining attributes?

Intro to the RELATIONAL model

Another KEY point

- do you see that, given these definitions, it MUST be the case that, for a relation scheme (e.g., $R(A_1, A_2, a_3, a_4, a_5)$), the primary key attributes MUST functionally determine the remaining attributes?
- (that is, $(A_1, A_2) \rightarrow (a_3, a_4, a_5)$)

Question

Consider the following relation, expressed in relation structure form:

Relly(ATT_A, ATT_B, att_c, att_d, att_e)

Assume this is a "good", real relation. Based on our relational definitions so far, which of the following functional dependencies must be the case?

- 1) $\text{att_a} \rightarrow \text{att_c}$
- 2) $\text{att_c} \rightarrow (\text{att_a}, \text{att_b})$
- 3) $(\text{att_a}, \text{att_b}) \rightarrow \text{att_e}$
- 4) $(\text{att_c}, \text{att_d}, \text{att_e}) \rightarrow (\text{att_a}, \text{att_b})$
- 5) $\text{att_b} \rightarrow (\text{att_c}, \text{att_d}, \text{att_e})$

Intro to the RELATIONAL model

ASIDE: the concept of physical keys/indexes

Intro to the RELATIONAL model

ASIDE: the concept of physical keys/indexes

- some DBMSs support identifying certain attributes in a table and it will "add" support to more quickly access those attributes;

Intro to the RELATIONAL model

ASIDE: the concept of physical keys/indexes

- some DBMSs support identifying certain attributes in a table and it will "add" support to more quickly access those attributes;
- Consider this relation:
- advisor = (advisor_id, advisor_lname, advisor_fname, advisor_phone)

Intro to the RELATIONAL model

ASIDE: the concept of physical keys/indexes

- some DBMSs support identifying certain attributes in a table and it will "add" support to more quickly access those attributes;
- these are NOT the same as primary keys! they are just for efficiency of certain searches...

Intro to the RELATIONAL model

ONE more important point... for future use:

- we've just said that every primary key must be the determinant of a functional dependency involving that relation, such that the primary key determines the rest of the attributes in that relation;

Intro to the RELATIONAL model

ONE more important point... for future use:

- we've just said that every primary key must be the determinant of a functional dependency involving that relation, such that the primary key determines the rest of the attributes in that relation;
 - BUT -- that is NOT to say that EVERY determinant in a relation is a primary key!

Intro to the RELATIONAL model

ONE more important point... for future use:

- we've just said that every primary key must be the determinant of a functional dependency involving that relation, such that the primary key determines the rest of the attributes in that relation;
 - BUT -- that is NOT to say that EVERY determinant in a relation is a primary key!
- Consider this relation:
 - student(STU_ID, stu_ssn, stu_lname, stu_fname, stu_email)
 - ...but `stu_ssn` is not the candidate key I chose for student's primary key.

Intro to the RELATIONAL model

- We've just said that every primary key must be the determinant of a functional dependency involving that relation, such that the primary key determines the rest of the attributes in that relation;
 - BUT -- that is NOT to say that EVERY determinant in a relation is a primary key!
- Consider this relation:
- ClassTaker(CLASSTAKER_ID, ctaker_Iname, ctaker_fname, ctaker_phone, advisor_id, advisor_Iname, advisor_phone)
 - Here, it IS the case that
 - $(\text{classtaker_id}) \rightarrow (\text{ctaker_Iname}, \text{ctaker_fname}, \text{ctaker_phone}, \text{advisor_id}, \text{advisor_Iname}, \text{advisor_phone})$

Intro to the RELATIONAL model

- BUT -- that is NOT to say that EVERY determinant in a relation is a primary key!
- ClassTaker(CLASSTAKER_ID, ctaker_lname, ctaker_fname, ctaker_phone, advisor_id, advisor_lname, advisor_phone)
 - Here, it IS the case that
 - $(\text{classtaker_id}) \rightarrow (\text{ctaker_lname}, \text{ctaker_fname}, \text{ctaker_phone}, \text{advisor_id}, \text{advisor_lname}, \text{advisor_phone})$
 - BUT it is also the case that
 - $(\text{advisor_id}) \rightarrow (\text{advisor_lname}, \text{advisor_phone})$
 - ...so in this case, advisor_id IS a determinant of a functional dependency, BUT it is not the primary key or even a candidate key for the relation ClassTaker

Intro to the RELATIONAL model

what, then, is a FOREIGN key?

- (broader/theoretical definition)
- a foreign key is a set of attributes within one relation that is also a CANDIDATE key of another relation

Intro to the RELATIONAL model

what, then, is a FOREIGN key?

- (broader/theoretical definition)
- a foreign key is a set of attributes within one relation that is also a CANDIDATE key of another relation
- BUT that's hard for DBMSs to determine automatically! SO in practice, many/most DBMSs require a foreign key to be a set of attributes within one relation that is also a PRIMARY key of another relation

Intro to the RELATIONAL model

what, then, is a FOREIGN key?

- (broader/theoretical definition)
- a foreign key is a set of attributes within one relation that is also a CANDIDATE key of another relation
- BUT that's hard for DBMSs to determine automatically! SO in practice, many/most DBMSs require a foreign key to be a set of attributes within one relation that is also a PRIMARY key of another relation
- these let us relate the tuples/rows in different relations in USEFUL ways... 8-)

Intro to the RELATIONAL model

some INTEGRITY definitions:

- referential integrity is when the foreign key attributes in a table are either null or also values in the referenced table
- ...then your data has referential integrity;

Intro to the RELATIONAL model

some INTEGRITY definitions:

- referential integrity is when the foreign key attributes in a table are either null or also values in the referenced table
- ...then your data has referential integrity;
- The foreign key can be null in certain cases because not all relationships between tables are required or mandatory. In those situations, a null value indicates the absence of a relationship.

Intro to the RELATIONAL model

some INTEGRITY definitions:

- referential integrity is when the foreign key attributes in a table are either null or also values in the referenced table
- ...then your data has referential integrity;
- ...DBMSs frequently do support referential integrity! Oracle does!

Intro to the RELATIONAL model

some INTEGRITY definitions:

- referential integrity is when the foreign key attributes in a table are either null or also values in the referenced table
- ...then your data has referential integrity;
- ...DBMSs frequently do support referential integrity! Oracle does!
- ...you make a set of attribute/attributes a foreign key, and Oracle enforces referential integrity

Question

You have table parts, with relation structure:

parts(PART_NUM, part_name, part_quant, part_price)

and table part_orders, with relation structure:

part_orders(ORD_NUM, part_num, ord_quant)

foreign key (part_num) references part

You try to insert an order with a part number that does not exist in the parts table, and the DBMS does not allow it. Which kind of integrity is the DBMS enforcing in this case?

- 1) entity integrity
- 2) domain integrity
- 3) referential integrity

Intro to the RELATIONAL model

some INTEGRITY definitions:

- entity integrity is when in a relation, no attribute of the primary key can be null AND no two rows can have the SAME value for the primary key

Intro to the RELATIONAL model

some INTEGRITY definitions:

- entity integrity is when in a relation, no attribute of the primary key can be null AND no two rows can have the SAME value for the primary key
- ...then your data has entity integrity;

Intro to the RELATIONAL model

some INTEGRITY definitions:

- entity integrity is when in a relation, no attribute of the primary key can be null AND no two rows can have the SAME value for the primary key
- ...then your data has entity integrity;
- ...DBMSs frequently do support entity integrity! Oracle does!
- ...you make a set of attribute/attributes a primary key, and Oracle enforces entity integrity

Question

You have table parts, with relation structure:

parts(PART_NUM, part_name, part_quant, part_price)

and table part_orders, with relation structure:

part_orders(ORD_NUM, part_num, ord_quant)

foreign key (part_num) references part

You try to insert a part with no part number, and the DBMS does not allow it. Which kind of integrity is the DBMS enforcing in this case?

- 1) entity integrity
- 2) domain integrity
- 3) referential integrity

Intro to the RELATIONAL model
some CLASSIC relational algebra operations

Intro to the RELATIONAL model

some CLASSIC relational algebra operations

- in relational algebra, the relational operations manipulate relations to form new relations!
- and, the results of a relational operation on relations are also relations
- (and because relations are sets, set operations work on them also, but we'll get to most of those later)

Intro to the RELATIONAL model

some CLASSIC relational algebra operations

- these are just a useful subset of the relational algebra operations that there are:
 - Selection
 - Projection
 - (Cartesian) product
 - natural join
 - equi-join

Intro to the RELATIONAL model

- Selection
 - when you select JUST specified tuples/rows from a relation

Intro to the RELATIONAL model

- Selection

For example, say that you have a **Student** relation as follows:

STUDENT table:

Stud_ID	Stud_LName	Stud_Major	Stud_Grade	Stud_Age_Level
123	Jones	History	JR	21
158	Parks	Math	GR	26
105	Anderson	Management	SR	27
271	Smith	History	JR	19

Then the result of the **selection** operation on the **Student** table of rows in which **Stud_Age < 25** would be:

STUDENT WHERE Stud_Age < 25

Stud_ID	Stud_LName	Stud_Major	Stud_Grade	Stud_Age_Level
123	Jones	History	JR	21
271	Smith	History	JR	19

Intro to the RELATIONAL model

- Projection
 - when you select JUST specified attributes from a relation (and make sure the result has no duplicate rows)

Intro to the RELATIONAL model

- **Projection** For example, say that you have a **Student** relation as follows:

STUDENT table:

Stud_ID	Stud_LName	Stud_Major	Stud_Grade_Level	Stud_Age
123	Jones	History	JR	21
158	Parks	Math	GR	26
105	Anderson	Management	SR	27
271	Smith	History	JR	19

Then the result of the **projection** operation of the **Student_Major** and **Stud_Grade_Level** attributes of the **Student** table would be:

Stud_Major	Stud_Grade_Level
History	JR
Math	GR
Management	SR

Question

STUDENT table:

Stud_ID	Stud_LName	Stud_Major	Stud_Grade_Level	Stud_Age
123	Jones	History	JR	21
158	Parks	Math	GR	26
105	Anderson	Management	SR	27
271	Smith	History	JR	19

Consider the displayed Student table. What would be the result of projecting the stud_major and stud_name column from the selection of rows in which stud_grade_level is GR?

(Assume each includes the appropriate column names/headers.)

1. Math Parks
2. Parks Math GR
3. 158 Parks Math GR 26
4. stud_major stud_name

Intro to the RELATIONAL model

- (Cartesian) product
 - the set of ALL pairs of ONE element from one set and ONE element from another set
 - ...but when relations are involved, the sets are sets of tuples!

Intro to the RELATIONAL model

- (Cartesian) product

For example, say that you have a **Price** relation as follows:

Prod_Code	Price
AA	5.99
BB	22.75

And, say that you have a **Location** table as follows:

Store	Aisle	Shelf
23	W	5
24	K	9
25	Z	6

Then the **Cartesian product** of Price and Location would be:

Prod_Code	Price	Store	Aisle	Shelf
AA	5.99	23	W	5
AA	5.99	24	K	9
AA	5.99	25	Z	6
BB	22.75	23	W	5
BB	22.75	24	K	9
BB	22.75	25	Z	6

Question

Assume that you have a table **Empl** with 20 rows. Which of the following is the case if you take the Cartesian product:

Empl x Empl

1. The resulting relation will have 20 rows.
2. The resulting relation will have 40 rows.
3. The resulting relation will have 400 rows.
4. The resulting relation will have no rows.

Intro to the RELATIONAL model

- equi-join, natural join
 - equi-join: when you take the Cartesian product of two relations, THEN you SELECT just the tuples in that product that meet some condition (usually an attribute value in one tuple being the same as an attribute value in another tuple from another relation)
 - MOST typically: you select the tuples in the Cartesian product that have common values for the attribute(s) serving as the foreign key relating those two tables

Intro to the RELATIONAL model

- equi-join, natural join
 - equi-join: when you take the Cartesian product of two relations, THEN you SELECT just the tuples in that product that meet some condition (usually an attribute value in one tuple being the same as an attribute value in another tuple from another relation)
 - MOST typically: you select the tuples in the Cartesian product that have common values for the attribute(s) serving as the foreign key relating those two tables
 - natural join: is the above, THEN you project ALL but one of the attributes/columns you joined on;

Intro to the RELATIONAL model

- equi-join, natural join

For example, say that you have **Student** and **Enrollment** tables as follows:

STUDENT table:

Stud_ID	Stud_LName	Stud_Major	Stud_Grade_Level	Stud_Age
---------	------------	------------	------------------	----------

123	Jones	History	JR	21
158	Parks	Math	GR	26
105	Anderson	Management	SR	27
271	Smith	History	JR	19

ENROLLMENT table:

Stud_ID	Class_Name	Position_Num
---------	------------	--------------

123	H350	1
105	BA490	3
123	BA490	7

Say that you wish to compute the **equi-join** and **natural join** of these tables based on the join condition (**Student.stud_id** = **Enrollment.stud_id**).

RELATIONAL model

- equi-join, natural join

First, compute the **Cartesian product** of these two tables:

Student. Stud_ID	Stud_ LName	Stud_ Major	Stud_ Grade_	Stud_ Age	Enrollment .Stud_ID	Class_ Name	Position _Num
Level							
123	Jones	History	JR	21	123	H350	1
123	Jones	History	JR	21	105	BA490	3
123	Jones	History	JR	21	123	BA490	7
158	Parks	Math	GR	26	123	H350	1
158	Parks	Math	GR	26	105	BA490	3
158	Parks	Math	GR	26	123	BA490	7
105	Anderson	Management	SR	27	123	H350	1
105	Anderson	Management	SR	27	105	BA490	3
105	Anderson	Management	SR	27	123	BA490	7
271	Smith	History	JR	19	123	H350	1
271	Smith	History	JR	19	105	BA490	3
271	Smith	History	JR	19	123	BA490	7

Second, perform a selection on this result of only those rows for which (**Student.stud_id = Enrollment.stud_id**):

Student. Stud_ID	Stud_ LName	Stud_ Major	Stud_ Grade_	Stud_ Age	Enrollment .Stud_ID	Class_ Name	Position _Num
Level							
123	Jones	History	JR	21	123	H350	1
123	Jones	History	JR	21	123	BA490	7
105	Anderson	Management	SR	27	105	BA490	3

This is the **equi-join** of these two tables on the join condition (**Student.stud_id = Enrollment.stud_id**).

Then, the **natural join** of these two tables on this join condition would include the **third** step of now projecting all of the columns in this result except for one of the "duplicate-contents" columns (and it doesn't matter which one of the two is omitted).

So, for example, the resulting **natural join** in this case could be:

Stud_ID	Stud_ LName	Stud_ Major	Stud_ Grade_	Stud_ Age	Class_ Name	Position_ Num
123	Jones	History	JR	21	H350	1
123	Jones	History	JR	21	BA490	7
105	Anderson	Management	SR	27	BA490	3

Question

Which has more columns -- the natural join of two tables, or the equi-join of two tables?

- 1) The natural join
- 2) The equi-join
- 3) Neither -- both always have the same number of columns
- 4) In general, one cannot say for sure.

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

Last login: Fri Sep 13 16:33:10 on ttys002

[di313@CS-S34217 ~ % ssh di313@nrs-projects-ssh.humboldt.edu

[di313@nrs-projects-ssh.humboldt.edu's password:

Activate the web console with: systemctl enable --now cockpit.socket

Last failed login: Fri Sep 13 16:45:07 PDT 2024 from 137.150.82.116 on ssh:notty

There was 1 failed login attempt since the last successful login.

Last login: Fri Sep 13 16:33:27 2024 from 137.150.82.116

[[di313@nrs-projects ~]\$ mkdir f24-325lect04-1

[[di313@nrs-projects ~]\$ chmod 700 f24-325lect04-1

[[di313@nrs-projects ~]\$ cd f24-325lect04-1

[di313@nrs-projects f24-325lect04-1]\$

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

Last login: Fri Sep 13 16:33:10 on ttys002

[di313@CS-S34217 ~ % ssh di313@nrs-projects-ssh.humboldt.edu

[di313@nrs-projects-ssh.humboldt.edu's password:

Activate the web console with: systemctl enable --now cockpit.socket

Last failed login: Fri Sep 13 16:45:07 PDT 2024 from 137.150.82.116 on ssh:notty

There was 1 failed login attempt since the last successful login.

Last login: Fri Sep 13 16:33:27 2024 from 137.150.82.116

[[di313@nrs-projects ~]\$ mkdir f24-325lect04-1

[[di313@nrs-projects ~]\$ chmod 700 f24-325lect04-1

[[di313@nrs-projects ~]\$ cd f24-325lect04-1

[di313@nrs-projects f24-325lect04-1]\$ vim 325lect04-1.sql

```
/* let's express relational operations in SQL */
```

```
/* pure selection! in SQL
```

```
select *  
from   desired_table  
where  desired_bool_expr;
```

```
*/
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
-- INSERT --
```

```
/* let's express relational operations in SQL */
```

```
-- get some example tables to play with
```

```
start set-up-ex-tbls.sql
```

```
/* pure selection! in SQL
```

```
select *
  from  desired_table
 where desired_bool_expr;
```

```
*/
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
-- INSERT --
```

The screenshot shows a dark-themed code editor window with a sidebar on the left containing various icons for file operations like copy, search, and refresh. The main area displays an SQL script named 'set-up-ex-tbls.sql'.

File Path: R: > Fall-2024 > CS 325 > Scripts > set-up-ex-tbls.sql

```
1 -----  
2 -- this file sets up example tables for query practice.  
3 --  
4 -- adapted from Oracle example tables used for SQL instruction  
5 --  
6 -- last modified: 2019-09-20  
7 -- by: Sharon Tuttle  
8 -----  
9  
10 -----  
11 -- create and populate table dept  
12  
13 drop table dept cascade constraints;  
14  
15 create table dept  
16 (dept_num    char(3),  
17 dept_name   varchar2(15) not null,  
18 dept_loc    varchar2(15) not null,  
19 primary key (dept_num)  
20 );  
21  
22 insert into dept  
23 values  
24 ('100', 'Accounting', 'New York');  
25  
26 insert into dept  
27 values  
28 ('200', 'Research', 'Dallas');  
29  
30 . . .
```

Bottom status bar: Restricted Mode, 0 ▲ 0, 0, Ln 1, Col 1, Spaces: 4, UTF-8, LF, MS SQL, 0 notifications.

The screenshot shows a dark-themed code editor window with a sidebar on the left containing various icons for file operations like copy, search, and refresh. The main area displays a SQL script named 'set-up-ex-tbls.sql'. The script creates a table 'empl' with columns for employee number, last name, job title, manager, hire date, salary, and department number. It includes primary key and foreign key constraints, and inserts data into the table.

```
R: > Fall-2024 > CS 325 > Scripts > set-up-ex-tbls.sql
43  -- create and populate table empl
44
45  drop table empl cascade constraints;
46
47  create table empl
48  (empl_num      char(4),
49   empl_last_name varchar2(15) not null,
50   job_title      varchar2(10),
51   mgr            char(4),
52   hiredate       date        not null,
53   salary          number(6,2),
54   commission     number(6,2),
55   dept_num       char(3),
56   primary key    (empl_num),
57   foreign key    (dept_num) references dept,
58   foreign key    (mgr)      references empl(empl_num));
59
60  insert into empl(empl_num, empl_last_name, job_title, hiredate,
61 | | | | salary, dept_num)
62 values
63 ('7839', 'King', 'President', '17-Nov-2011', 5000.00, '500');
64
65  insert into empl(empl_num, empl_last_name, job_title, mgr, hiredate,
66 | | | | salary, dept_num)
67 values
68 ('7566', 'Jones', 'Manager', '7839', '02-Apr-2012', 2975.00, '200');
69
70  insert into empl(empl_num, empl_last_name, job_title, mgr, hiredate,
71 | | | | salary, dept_num)
72 values
```

Bottom status bar: Restricted Mode, 0 ▲ 0, 0, Ln 1, Col 1, Spaces: 4, UTF-8, LF, MS SQL, 0

The screenshot shows a code editor interface with a dark theme. On the left, there is a vertical toolbar with various icons: a file icon, a search icon, a copy/paste icon, a refresh icon, a dropdown menu icon, a zoom icon, a settings gear icon, and a user profile icon.

The main area displays an SQL script named "set-up-ex-tbls.sql". The script is used to create and populate a table named "customer". It starts with dropping the table if it exists, then creating it with columns for customer ID, last name, first name, employee representative, street address, city, state, zip code, balance, and primary key. It includes foreign key constraints linking to an "empl" table. Finally, it inserts two rows of data into the "customer" table.

```
R: > Fall-2024 > CS 325 > Scripts > set-up-ex-tbls.sql

130  -----
131  -- create and populate table customer
132
133  drop table customer cascade constraints;
134
135  create table customer
136  (cust_id      char(6),
137   cust_lname    varchar2(20)  not null,
138   cust_fname    varchar2(15),
139   empl_rep     char(4),
140   cust_street   varchar2(30),
141   cust_city     varchar2(15),
142   cust_state    char(2),
143   cust_zip      varchar2(10),
144   cust_balance  number(7,2)   default 0.0,
145   primary key  (cust_id),
146   foreign key  (empl_rep) references empl(empl_num));
147
148  insert into customer
149  values
150  ('100001', 'Firstly', 'First', '7499', '1111 First Street', 'Fortuna', 'CA',
151  | '95520', 1111.11);
152
153  insert into customer
154  values
155  ('100002', 'Secondly', 'Second', '7654', '2222 Second Street',
156  | 'McKinleyville', 'CA', '95523', 222.20);
157
158  insert into customer(cust_id, cust_lname, cust_fname, empl_rep,
```

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

[di313@nrs-projects f24-325lect04-1]\$ sqlplus /

SQL*Plus: Release 19.0.0.0.0 - Production on Mon Sep 16 07:34:21 2024
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.

Last Successful login time: Mon Sep 16 2024 07:30:28 -07:00

Connected to:

Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0

SQL> @ 325lect04-1.sql

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

1 row created.

1 row created.

1 row created.

Table dropped.

Table created.

1 row created.

1 row created.

1 row created.

SQL> □

```
/* let's express relational operations in SQL */
```

```
-- get some example tables to play with
```

```
start set-up-ex-tbls.sql
```

```
/* pure selection! in SQL
```

```
select *
from   desired_table
where  desired_bool_expr;
```

```
*/
```

```
-- selection of empls whose salary is more than 1500
```

```
select *
from   empl
where  salary > 1500;~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
-- INSERT --
```

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

1 row created.

1 row created.

1 row created.

Table dropped.

Table created.

1 row created.

1 row created.

1 row created.

SQL> @ 325lect04-1.sql

1 row created.

1 row created.

1 row created.

EMPL	EMPL_LAST_NAME	JOB_TITLE	MGR	HIREDATE	SALARY	COMMISSION	DEP
7839	King	President		17-NOV-11	5000		500
7566	Jones	Manager	7839	02-APR-12	2975		200
7698	Blake	Manager	7839	01-MAY-13	2850		300
7782	Raimi	Manager	7839	09-JUN-12	2450		100
7902	Ford	Analyst	7566	03-DEC-12	3000		200
7499	Michaels	Sales	7698	20-FEB-18	1600	300	300
7788	Scott	Analyst	7566	09-NOV-18	3000		200

7 rows selected.

SQL> □

```
-- selection of empls whose salary is more than 1500
```

```
select *  
from   empl  
where  salary > 1500;
```

```
/* pure projection!
```

```
  select distinct attrib1, attrib2, ...  
  from desired_tbl;  
*/
```

```
-- projection of dept_num and job_title
```

```
-- selection of empls whose salary is more than 1500
```

```
select *  
from   empl  
where  salary > 1500;
```

```
/* pure projection!
```

```
  select distinct attrib1, attrib2, ...  
  from desired_tbl;  
*/
```

```
-- projection of dept_num and job_title
```

```
select distinct dept_num, job_title  
from   empl;
```

```
□
```

1 row created.

1 row created.

1 row created.

EMPL	EMPL_LAST_NAME	JOB_TITLE	MGR	HIREDATE	SALARY	COMMISSION	DEP
7839	King	President		17-NOV-11	5000	500	
7566	Jones	Manager	7839	02-APR-12	2975	200	
7698	Blake	Manager	7839	01-MAY-13	2850	300	
7782	Raimi	Manager	7839	09-JUN-12	2450	100	
7902	Ford	Analyst	7566	03-DEC-12	3000	200	
7499	Michaels	Sales	7698	20-FEB-18	1600	300	300
7788	Scott	Analyst	7566	09-NOV-18	3000		200

7 rows selected.

SQL> @ 325lect04-1.sql

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

7782	Raimi	Manager	7839	09-JUN-12	2450	100
7902	Ford	Analyst	7566	03-DEC-12	3000	200
7499	Michaels	Sales	7698	20-FEB-18	1600	300 300
7788	Scott	Analyst	7566	09-NOV-18	3000	200

7 rows selected.

DEP JOB_TITLE

200	Analyst
300	Clerk
500	President
100	Manager
200	Clerk
300	Manager
200	Manager
100	Clerk
300	Sales
400	Clerk

10 rows selected.

SQL>

Question

product_id	product_name	price	category
1	Laptop	1000	Electronics
2	Smartphone	800	Electronics
3	Chair	200	Furniture
4	Desk	400	Furniture
5	TV	1200	Electronics

Which SQL query will return only the product_name and category for products priced above 500, without duplicates?

- 1) SELECT product_name, category FROM products WHERE price > 500;
- 2) SELECT DISTINCT product_name, category FROM products WHERE price > 500;
- 3) SELECT product_name, category FROM products WHERE price < 500;
- 4) SELECT DISTINCT * FROM products WHERE price > 500;

*/

```
-- projection of dept_num and job_title
```

```
select distinct dept_num, job_title  
from    empl;
```

```
/* pure Cartesian product!
```

```
select *  
from    tbl1, tb2;
```

*/

```
-- Cartesian product of empl and dept (ugly!)
```

```
□  
select *  
from    empl, dept;
```

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

7782	Raimi	Manager	7839	09-JUN-12	2450	100
7902	Ford	Analyst	7566	03-DEC-12	3000	200
7499	Michaels	Sales	7698	20-FEB-18	1600	300 300
7788	Scott	Analyst	7566	09-NOV-18	3000	200

7 rows selected.

DEP JOB_TITLE

200	Analyst
300	Clerk
500	President
100	Manager
200	Clerk
300	Manager
200	Manager
100	Clerk
300	Sales
400	Clerk

10 rows selected.

SQL> @ 325lect04-1.sql

```
DEPT_NAME      DEPT_LOC
```

```
7844 Turner      Sales      7698 08-SEP-19      1500      0 300 500
Management      New York
```

```
7876 Adams      Clerk       7788 23-SEP-18      1100      400 500
Management      New York
```

```
7900 James      Clerk       7698 03-DEC-17      950       300 500
Management      New York
```

```
EMPL  EMPL_LAST_NAME  JOB_TITLE  MGR  HIREDATE      SALARY  COMMISSION  DEP  DEP
```

```
DEPT_NAME      DEPT_LOC
```

```
7934 Miller      Clerk       7782 23-JAN-16      1300      100 500
Management      New York
```

70 rows selected.

SQL> □

```
from    tbl1, tb2;  
  
*/  
  
-- Cartesian product of empl and dept (ugly!)  
  
select *  
from    empl, dept;  
  
/* ONE form of pure equi-join!  
  
select *  
from    tbl1, tbl2  
where  tbl1.attrib = tbl2.attrib;  
  
*/  
  
select *  
from    empl, dept  
where  empl.dept_num = dept.dept_num;□
```

di313 — di313@nrs-projects:~/f24-325lect04-1— ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

DEPT_NAME DEPT_LOC

7844	Turner Management	Sales New York	7698	08-SEP-19	1500	0	300	500
------	----------------------	-------------------	------	-----------	------	---	-----	-----

7876	Adams Management	Clerk New York	7788	23-SEP-18	1100	400	500
------	---------------------	-------------------	------	-----------	------	-----	-----

7900	James Management	Clerk New York	7698	03-DEC-17	950	300	500
------	---------------------	-------------------	------	-----------	-----	-----	-----

EMPL EMPL_LAST_NAME JOB_TITLE MGR HIREDATE SALARY COMMISSION DEP DEP

DEPT_NAME DEPT_LOC

7934	Miller Management	Clerk New York	7782	23-JAN-16	1300	100	500
------	----------------------	-------------------	------	-----------	------	-----	-----

70 rows selected.

SQL> @ 325lect04-1.sql

7788 Scott Research	Analyst Dallas	7566 09-NOV-18	3000	200 200
7844 Turner Sales	Sales Chicago	7698 08-SEP-19	1500	0 300 300
7876 Adams Operations	Clerk Boston	7788 23-SEP-18	1100	400 400

EMPL	EMPL_LAST_NAME	JOB_TITLE	MGR	HIREDATE	SALARY	COMMISSION	DEP	DEP
DEPT_NAME	DEPT_LOC							
7900 Sales	James	Clerk	7698	03-DEC-17	950		300	300
		Chicago						
7934 Accounting	Miller	Clerk	7782	23-JAN-16	1300		100	100
		New York						

14 rows selected.

SQL> □

DEPT_NAME	DEPT_LOC				
7900 James Sales	Clerk Chicago	7698	03-DEC-17	950	300 300
7934 Miller Accounting	Clerk New York	7782	23-JAN-16	1300	100 100

14 rows selected.

[SQL]> describe empl

Name	Null?	Type
EMPL_NUM	NOT NULL	CHAR(4)
EMPL_LAST_NAME	NOT NULL	VARCHAR2(15)
JOB_TITLE		VARCHAR2(10)
MGR		CHAR(4)
HIREDATE	NOT NULL	DATE
SALARY		NUMBER(6,2)
COMMISSION		NUMBER(6,2)
DEPT_NUM		CHAR(3)

SQL>

14 rows selected.

[SQL]> describe empl

Name	Null?	Type
EMPL_NUM	NOT NULL	CHAR(4)
EMPL_LAST_NAME	NOT NULL	VARCHAR2(15)
JOB_TITLE		VARCHAR2(10)
MGR		CHAR(4)
HIREDATE	NOT NULL	DATE
SALARY		NUMBER(6,2)
COMMISSION		NUMBER(6,2)
DEPT_NUM		CHAR(3)

[SQL]> describe dept

Name	Null?	Type
DEPT_NUM	NOT NULL	CHAR(3)
DEPT_NAME	NOT NULL	VARCHAR2(15)
DEPT_LOC	NOT NULL	VARCHAR2(15)

SQL>

*/

```
select *  
from   empl, dept  
where  empl.dept_num = dept.dept_num;
```

/* natural join is just... tedious!

```
select <all attrs of tbl1 and tbl2 except ONE of tbl1's or  
      tbl2's join_attrib>  
from   tbl1, tbl2  
where  tbl1.join_attrib = tbl2.join_attrib;
```

*/

```
select empl_num, empl_last_name, job_title, mgr, hiredate, salary,  
       commission, empl.dept_num, dept_name, dept_loc  
from   empl, dept  
where  empl.dept_num = dept.dept_num;
```

□

14 rows selected.

[SQL] > describe empl

Name	Null?	Type
EMPL_NUM	NOT NULL	CHAR(4)
EMPL_LAST_NAME	NOT NULL	VARCHAR2(15)
JOB_TITLE		VARCHAR2(10)
MGR		CHAR(4)
HIREDATE	NOT NULL	DATE
SALARY		NUMBER(6,2)
COMMISSION		NUMBER(6,2)
DEPT_NUM		CHAR(3)

[SQL] > describe dept

Name	Null?	Type
DEPT_NUM	NOT NULL	CHAR(3)
DEPT_NAME	NOT NULL	VARCHAR2(15)
DEPT_LOC	NOT NULL	VARCHAR2(15)

[SQL] > @ 325lect04-1.sql

di313 — di313@nrs-projects:~/f24-325lect04-1 — ssh di313@nrs-projects-ssh.humboldt.edu — 80x24

7788	Scott Research	Analyst Dallas	7566	09-NOV-18	3000	200
7844	Turner Sales	Sales Chicago	7698	08-SEP-19	1500	0 300
7876	Adams Operations	Clerk Boston	7788	23-SEP-18	1100	400

EMPL	EMPL_LAST_NAME	JOB_TITLE	MGR	HIREDATE	SALARY	COMMISSION	DEP
DEPT_NAME	DEPT_LOC						
7900	James Sales	Clerk Chicago	7698	03-DEC-17	950	300	
7934	Miller Accounting	Clerk New York	7782	23-JAN-16	1300	100	

14 rows selected.

SQL> □

Question

Which of the following are relational operations from relational algebra?

- 1) create table, drop table
- 2) selection, projection, natural join
- 3) insert, select
- 4) spool, describe